

## Traditional distributed systems

- Paxos
- raft

trusted environment

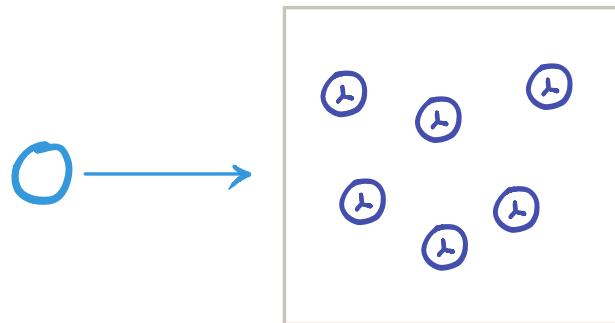
## Blockchain

- PoW (Proof of Work):
- PoS (Proof of Stake):
- DPoS (Delegated Proof of Stake):
- PoI (Proof of Importance):
- PoD (Proof of Devotion)
- PBFT (Practical Byzantine Fault Tolerance)
- FPA (Federated Byzantine Agreement)
- Hybrid PoW/PoS
- Proof-of-DDOS

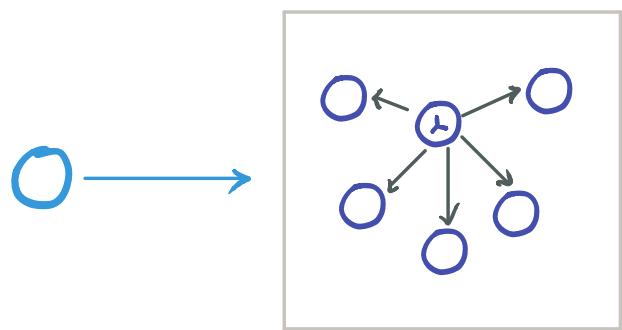
untrusted

Purpose: decrease the incentive for people to misbehave.

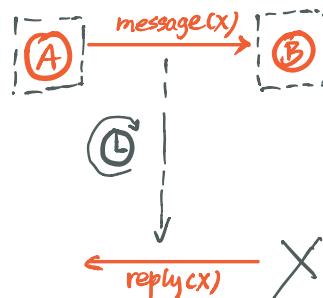
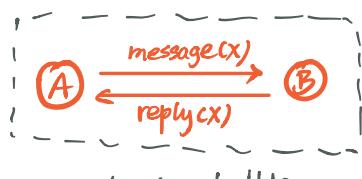
- unfeasible to attack

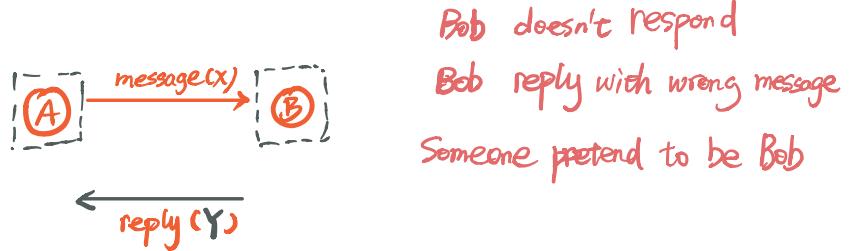


*Symmetric consensus*  
Any server can respond



*Asymmetric consensus*  
A leader is elected to issue command





Bob doesn't respond  
Bob reply with wrong message  
Someone pretend to be Bob

Blockchain - proof based , slow  
Paxos - hard to implement and understand

# Paxos

**Phase 1.** (a) A proposer selects a proposal number  $n$  and sends a *prepare* request with number  $n$  to a majority of acceptors.

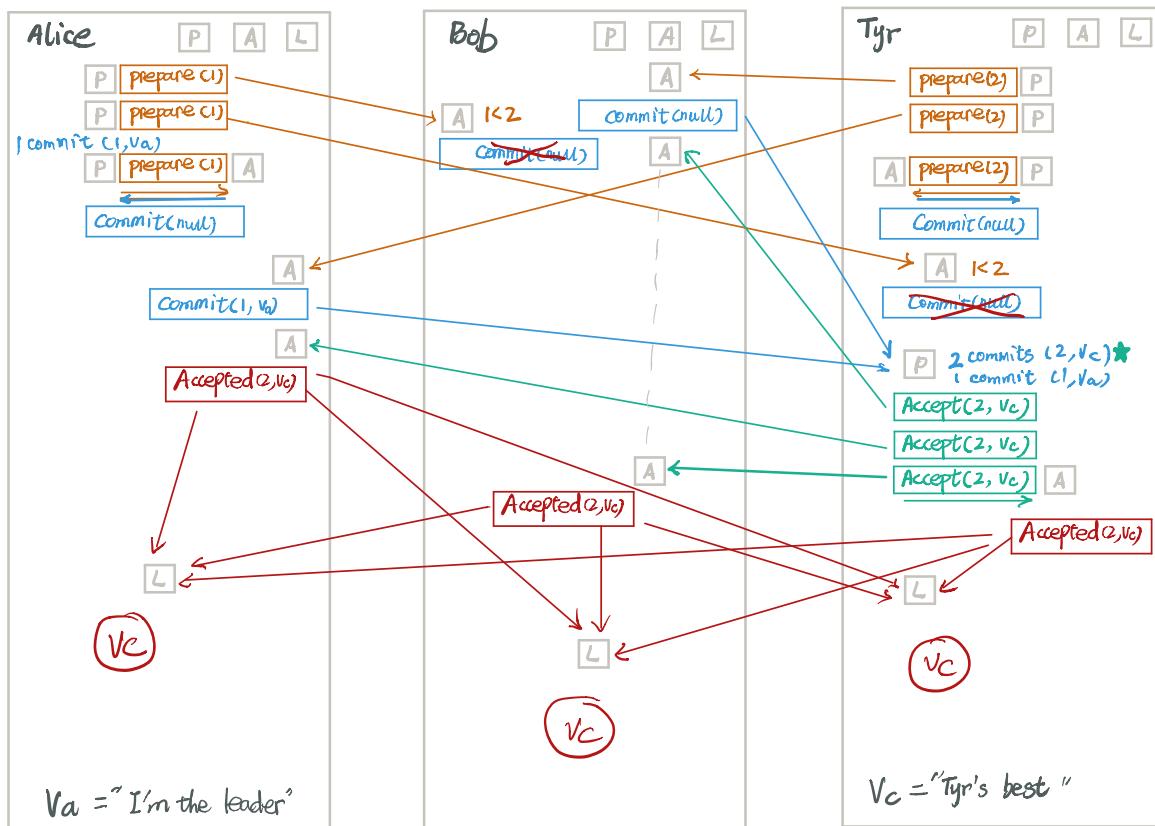
(b) If an acceptor receives a *prepare* request with number  $n$  greater than that of any *prepare* request to which it has already responded, then it responds to the request with a promise not to accept any more proposals numbered less than  $n$  and with the highest-numbered proposal (if any) that it has accepted.

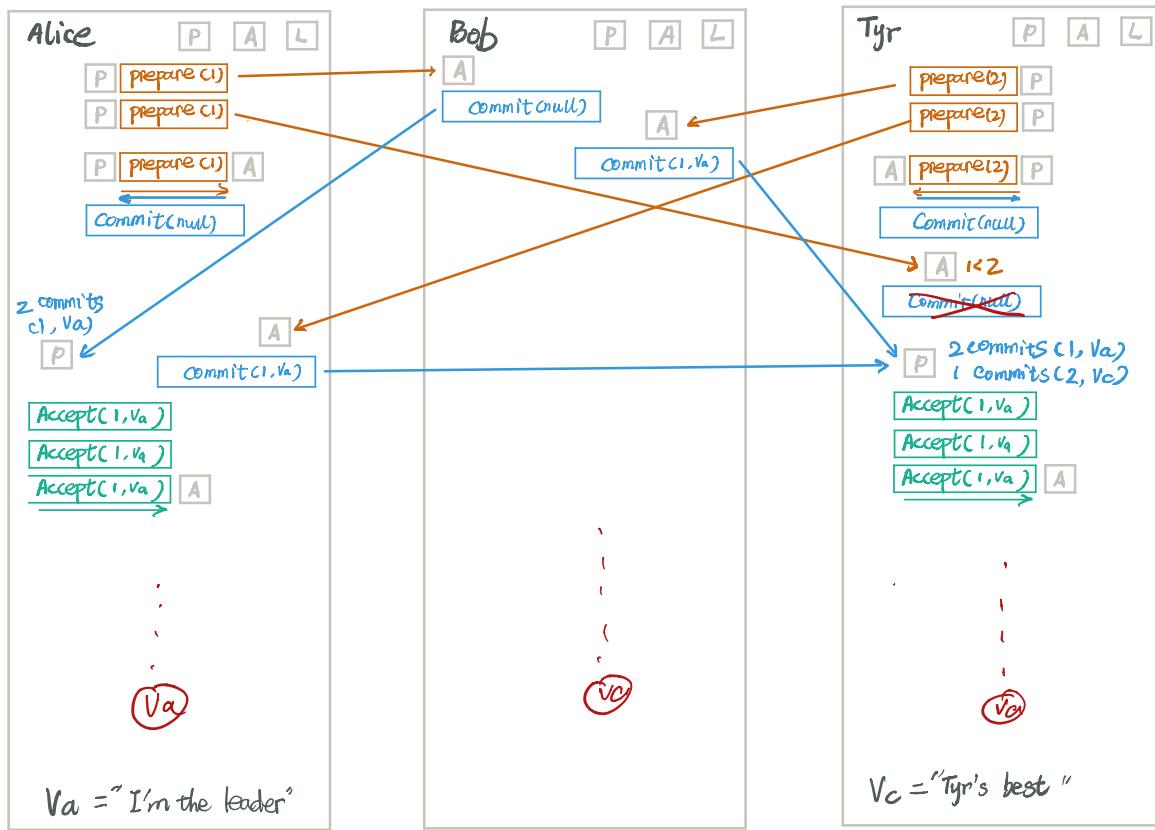
**Phase 2.** (a) If the proposer receives a response to its *prepare* requests (numbered  $n$ ) from a majority of acceptors, then it sends an *accept* request to each of those acceptors for a proposal numbered  $n$  with a value  $v$ , where  $v$  is the value of the highest-numbered proposal among the responses, or is any value if the responses reported no proposals.

(b) If an acceptor receives an *accept* request for a proposal numbered  $n$ , it accepts the proposal unless it has already responded to a *prepare* request having a number greater than  $n$ .

## Phase 3.

To learn that a value has been chosen, a learner must find out that a proposal has been accepted by a majority of acceptors. The obvious algorithm is to have each acceptor, whenever it accepts a proposal, respond to all learners, sending them the proposal.





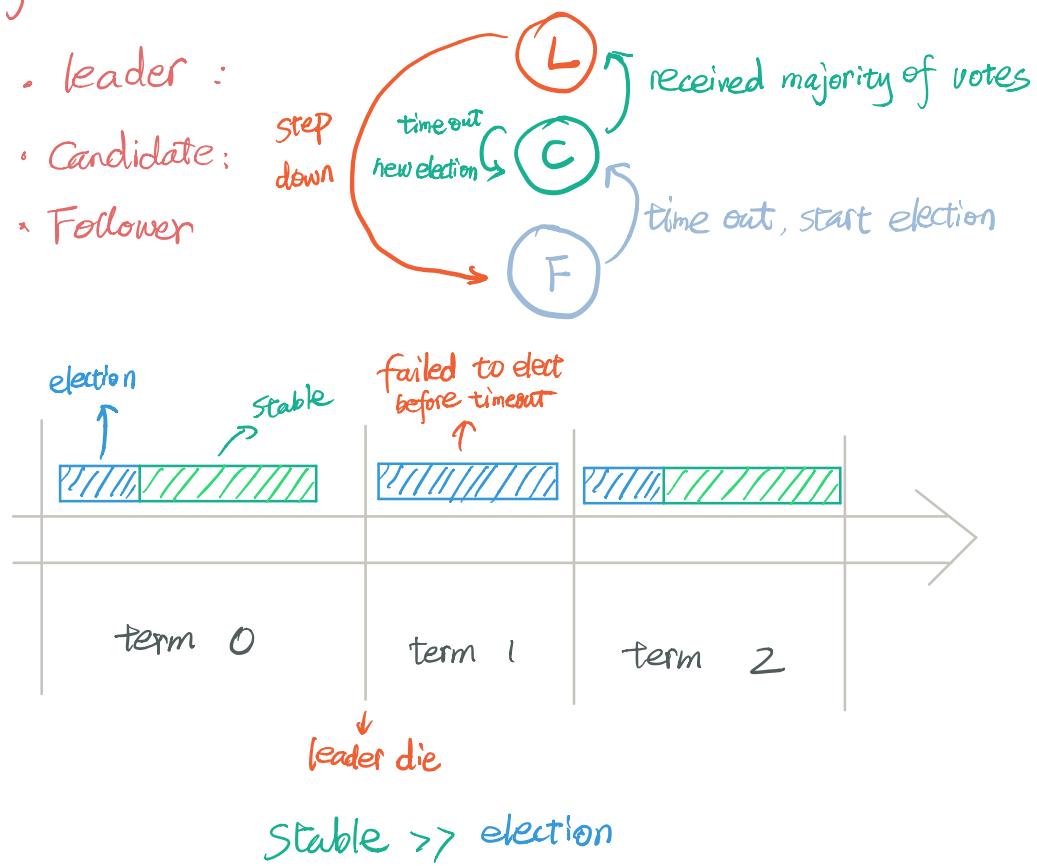
# Raft

Goal: Have and available replicated state machine.

{ available: provide meaningful response for a log  
replicated: makes other nodes aware of its log  
state machine: follow the log to arrive at the same state

## Asymmetric

- leader :
- Candidate:
- Follower



## Election

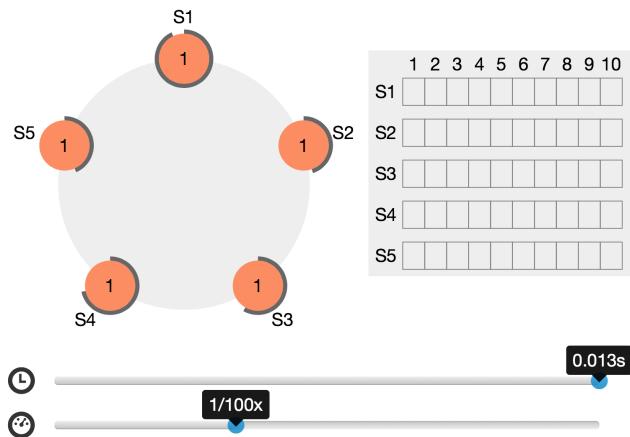
1. Increment current term
2. vote for self
3. request votes

properties {  
Safety : at most one winner per term.  
liveness: Someone eventually wins.

outcome {  
I'm leader: receive majority votes  
I'm not : receive valid heart beat  
no leader: election timeout

### Raft Visualization

Here's a Raft cluster running in your browser. You can interact with it to see Raft in action. Five servers are shown on the left, and their logs are shown on the right. We hope to create a screencast soon to explain what's going on. This visualization (RaftScope) is still pretty rough around the edges; pull requests would be very welcome.

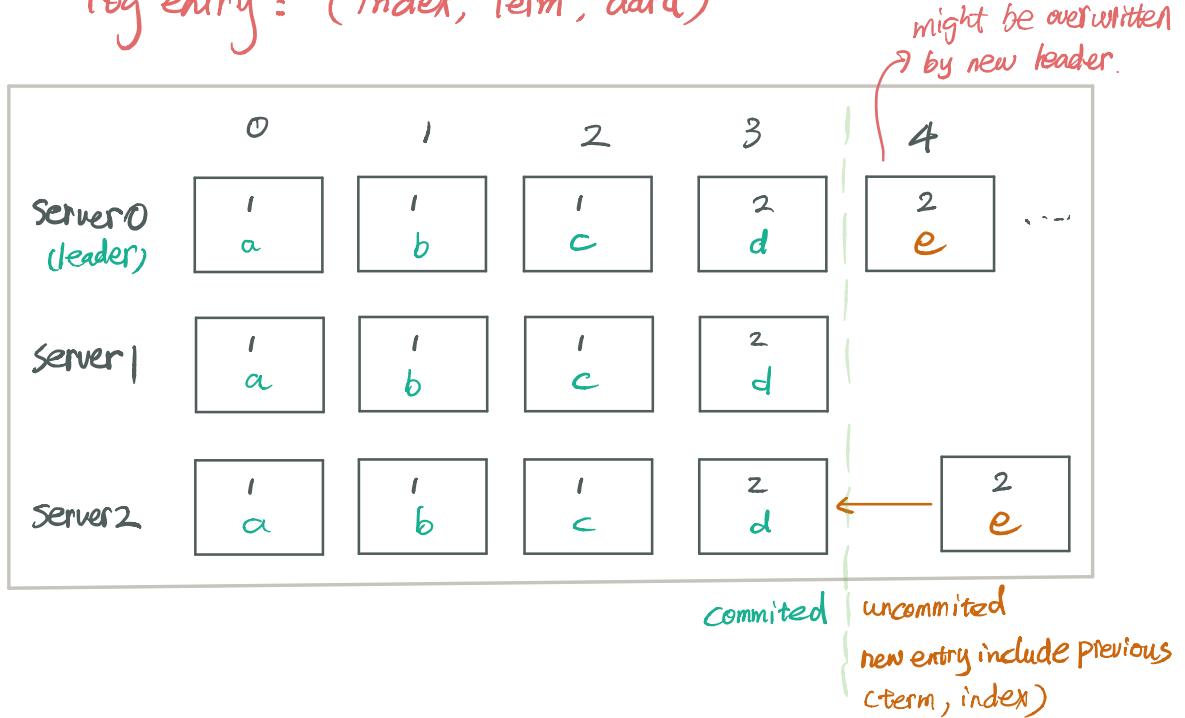


[The Secret Lives of Data](#) is a different visualization of Raft. It's more guided and less interactive, so it may be a gentler starting point.

## Log Replication

↳ exchange information

log entry : (index, term, data)



log replication

log safety : committed entries in the logs of all future leaders.  
log coherency : deny votes if candidate has "less complete" log  
log integrity : an entry is committed if 1) it's majority stored;  
2) at least one entry from leader's term is also majority stored.

who's using : Cloud Foundry, Kubernetes, etcd, docker Swarm, Consul

not vote { if your term is larger  
if your index is better  
 $(t_v > t_c)$  or  $((t_v = t_c) \text{ and } (i_v > i_c))$

# PoW

Pros:

- attack is hard, unfeasible
- the coin is precious given that so much has been spent.

Cons:

- amateur can't compete with professional miners. (CPU vs ASIC/GPU)
- Matthew Effect.

## PoS

Pros:

- Attackers have no incentive to attack
- Inexpensive, low cost.

Cons:

- Matthew Law

## PoI (Proof of Importance)

Not only reward those with large balance, but also take into account how much they transact to others and who they transact with.

Pros:

- Good incentive system. The more you use, the better your importance (trust score) is.